

Architecture and Requirements for Observability, C&I of Network Management Agents


An architecture and requirements briefing for supervising autonomous AI agents deployed for network operations

Qiufang Ma (Presenting)
Qin Wu

Scope

- Defines the architecture and requirements between an agent and its management plane
- Solution-neutral
 - Protocols, data model, or API choice is intentionally left for future work.
- Companion Problem Statement
 - I-D. wnd-opsawg-icon-ps details the operational gaps motivating these requirements

A Reference Architecture



Human Oversight - the top-level authority

- Agent runtime monitoring
- Policy & constraint injection
- Emergency intervention trigger
- Escalation request handling
- Post-execution feedback

Agent Management Plane – centralized layer for agent assurance

- Agent runtime telemetry data collection
- Control policy validation and distribution
- Agent runtime intervention instruction (e.g., kill switch, pause/resume, or rollback)

Agent Execution Plane – agent runtime environment

- LLM inference
- Policy-compliant network operation task execution
- Execution result reporting

Deriving Requirements from Problem Statements- Agent Observability

See what agents did, why they did it, and how they're performing.

Problem Statement Gap I-D. wnd-opsawg-icon-ps	Essential Requirement I-D. mcw-opsawg-icon-requirements	Requirement Description
3.1.1 Limited Transparency in Planning and Decision-Making	OBS-1: Execution Trajectory Capture	The full path visibility, including the plan, the invoked tools/skills, generated configuration change, and the resulting operational state
5.1 OpenTelemetry only captures the execution process, not the operational motivation.	OBS-2: Reasoning Provenance Capture	Exposes why a decision was made — the operational signals (e.g., alarms, incidents, telemetry streams) and confidence level behind a specific network action.
5.1 lacks native support for observing metrics such as an agent's reasoning logic and internal confidence levels	OBS-3: Operational Metrics Collection	Tracks an agent's operational health and performance when performing specific network operations.
3.1.2. Ambiguity of Accountability Attribution	OBS-4: Auditability and Accountability	Immutable audit logging to allows the attribution of failures to responsible entities

Deriving Requirements from Problem Statements- Agent Control

Set a boundary for agent behaviors

Problem Statement Gap I-D. wnd-opsawg-icon-ps	Essential Requirement I-D. mcw-opsawg-icon-requirements	Description
3.2.1. Inadequacy of Deterministic Constraints	CTL-1: Intent Validation & Alignment	Requires the agent to check that what it is about to configure actually matches the operational intent it was given, before committing.
3.2.1. Inadequacy of Deterministic Constraints	CTL-2: Temporal Validity	Ties an agent's actions to a valid operating window.
3.2.2 Static IAM Limitation	CTL-3: Access & Permission	Limits each agent's operational access scope to specific network domains/areas, devices, protocols and tools as well as network configuration (e.g., datastores, modules, nodes, RPCs.)
3.2.2 Static IAM Limitation	CTL-4: Authorization & Approval	Specifies of some network operations as requiring a human decision before execution.
3.2.2. Static IAM Limitation	CTL-5: Dynamic Boundary Adaptation	Across multi-agent environments, and enable dynamic adjustment of operational bounds (e.g., tighten the agent's permissible access from RW to RO) based on the network's current operational state.

Deriving Requirements from Problem Statements- Agent Intervention

Emergency runtime interruption when the agent deviates from expected behaviors.

Problem Statement Gap I-D. wnd-opsawg-icon-ps	Essential Requirement I-D. mcw-opsawg-icon-requirements	Description
3.3.1 lack of human oversight Essential human oversight capabilities including execution interruption 5.5 absence of an interface that allows an authorized authority outside the framework or outside the agent application to signal intervention; current practice of intervention lacks soft redirect, scope restriction, checkpoint, task suspension, rollback, hard termination	INT-1: Execution Interruption	stop or redirect a running agent's runtime execution. Guarantees an agent can be paused/stopped from outside its own decision loop, even if the agent itself is unresponsive or stuck.
3.3.1. Lack of Human Oversight Essential human oversight capabilities including deterministic transaction rollback and Recovery	INT-2: Rollback and Recovery	Able to reverse the operations performed by the agent on the network based on different rollback granularities: workflow level, task level, and context level.
3.3.1. Lack of Human Oversight Essential human oversight capabilities including human escalation	INT-3: Escalation	Moves a decision the current agent can't or shouldn't make alone up to someone with broader authority, without losing operational context.
4.6 Intervention and Control is concerned with ensuring actions can be corrected or reversed when necessary.	INT-4: Correction	Allows a supervisor to fix what went wrong, e.g., the intent, the data, or the plan instead of only being able to accept or reject the agent's output.

Comments, Questions, Concerns?